

## COSC 3360/6310 - Operating Systems Fall 2019

### Programming Assignment 2 - Deadlock Avoidance with EDF plus SJF-Tie-Breaker Scheduling and EDZL plus LJF-Tie-Breaker Scheduling

**Due Date: Friday, November 1, 2019, 11:59pm CDT**

In this assignment, you will implement a deadlock avoidance algorithm as part of the Process Manager to avoid deadlocks in a Unix/Linux system. Part of the assignment requires the manipulation of Unix/Linux processes and part of it consists of simulation.

Both the deadlock-handling process and the processes requesting resources are real Unix/Linux processes created using `fork()`. However, you do not need to actually allocate any resource. The main process executes the Banker's algorithm. The resource-requesting processes are required to make requests by communicating with the deadlock-handling process with Unix/Linux shared memory controlled by Unix/Linux semaphores.

The deadlock-handling process chooses the next process with a resource request having the nearest absolute deadline to be serviced. Ties are broken in favor of the process with the shortest remaining execution time (Shortest Job First - SJF). After having one request serviced, a process has to allow another process to make a request before making another one, that is, another process with the nearest absolute deadline is chosen for service. A process can also release resources during its execution, and releases all resources held when it terminates.

Associated with each process is also a relative deadline (that is, the deadline from the current time instant). One time unit is needed for servicing each request (whether the resource is allocated or not), so the relative deadline for each process is decreased by 1 whether it is serviced or not. If the resource is allocated for a request, then the computation time of this process decreases by 1; otherwise, the computation time remains the same. If a process misses its deadline, keep servicing its requests but indicate the deadline miss in the output. A 'release' instruction also needs 1 unit of computation time just like a request.

A 'calculate' instruction does not use resources and its computation time is indicated in parentheses. A 'useresources' instruction simulates the use of allocated resources and its computation time is indicated in parentheses.

The input format is as follows:

```
m          /* number of resources */
n          /* number of processes */

available[1] = number of instances of resource 1
:
available[m] = number of instances of resource m

max[1,1] = maximum demand for resource 1 by process 1
:
max[n,m] = maximum demand for resource m by process n

process_1:
```

```

deadline_1      /* an integer, must be >= computation time */
computation_time_1 /* an integer, equal to number of requests and releases */
:               /* plus the parenthesized values in the calculate and */
:               /* userresources instructions. */
:
calculate(2);    /* calculate without using resources */
calculate(1);
request(0,1,0,...,2); /* request vector, m integers */
userresources(4);    /* use allocated resources */
:
release(0,1,0,...,1); /* release vector, m integers */
calculate(3);
:
request(1,0,3,...,1); /* request vector, m integers */
userresources(5);
:
end.

:

process_n:
deadline_n      /* an integer */
computation_time_n /* an integer, equal to number of requests and releases */
:               /* plus the parenthesized values in the calculate and */
:               /* userresources instructions. */
:
calculate(3);    /* calculate without using resources */
:
request(0,2,0,...,2); /* request vector, m integers */
userresources(2);    /* use allocated resources */
userresources(5);
userresources(3);
:
release(0,1,0,...,2); /* release vector, m integers */
calculate(4);
calculate(5);
:
request(1,0,3,...,1); /* request vector, m integers */
userresources(8);
calculate(3);
:
end.

```

For output, print the state of the system after servicing each request: the arrays available, allocation, need, and deadline misses, if any.

Next, let's try EDZL (Earliest Deadline Zero Laxity) with longest remaining execution time (Longest Job First - LJF) tie breaker in the second version of your algorithm. EDZL is a variant of EDF in which processes with zero laxity or negative laxities are given highest priority even though

their absolute deadlines are not the closest, and other processes are ordered by their absolute deadlines. Processes with zero or negative laxities are scheduled according to the LLF (Least Laxity First) scheduling strategy, thus the deadlock-handling process chooses the next process (among those with zero or negative laxities) with the smallest laxity to be serviced. Ties are broken in favor of the process with the longest remaining execution time. After having one request serviced, a process has to allow another process to make a request before making another one, that is, another process with the highest priority according to EDZL is chosen for service.

Therefore, this project has two runs corresponding to two different schedulers (EDF and EDZL). Keep executing processes until they finish even if they have already missed their deadlines. Which scheduling technique yields fewer deadline misses?